

Tree-based Models

Zhiyu Quan

University of Connecticut

20 March, 2018

Introduction and motivation

- Nonparametric approach – distribution free
- Interpret by visualizing the tree structure
- Robust to the outliers and missing data
- Partially solves multicollinearity
- Detect non-linear effects and interactions

- Regression trees can fit almost all the traditional statistical models
 - Least-squares, Logistic, Poisson, Proportional Hazards models
 - Quantile, Longitudinal, multiresponse

CART (classification and regression tree)

- Grows a large tree
- Prune the large tree

CART (Grows a large tree)

- Decision tree($T(\mathbf{x}, \Theta)$) algorithm partition the explanatory variable space into disjoint regions M regions R_1, R_2, \dots, R_M ; then assign a constant c_m in each region which is the estimated values \hat{y}_i of the response variable y_i in region R_m :

$$\hat{y}_i = f(y_i | \mathbf{x}_i, \Theta) = T(\mathbf{x}_i; \Theta) = \sum_{m=1}^M c_m \mathbf{1}_{R_m}(\mathbf{x}_i) \text{ where } \Theta = \{R_m, c_m\}_{m=1}^M$$

- Under the sum of square error(SSE), recursive binary splitting, first finds the single numerical explanatory variable X_j which best splits the data into two regions $R_1(j, s) = \{\mathbf{x}_i | X_{.j} < s\}$ and $R_2(j, s) = \{\mathbf{x}_i | X_{.j} > s\}$, i.e. For any j and s ,

$$\operatorname{argmin}_{j, s} \sum_{i: \mathbf{x}_i \in R_1(j, s)} (y_i - \hat{c}_{R_1(j, s)})^2 + \sum_{i: \mathbf{x}_i \in R_2(j, s)} (y_i - \hat{c}_{R_2(j, s)})^2$$

CART (Prune the large tree)

- cost-complexity pruning:

$$C_{\alpha}(T) = \sum_{m=1}^{|T|} \sum_{\mathbf{x}_i \in R_m} (y_i - \hat{c}_m)^2 + \alpha |T|$$

- The tuning parameter $\alpha \geq 0$ governs the tradeoff between tree size and its goodness of fit to the data.

CART (Missing data)

- The CART employs a series of “surrogate” splits
- Splits on alternate explanatory variables that substitute for the best explanatory variable which has a missing value
- Surrogate splits are also used to provide variable importance.

CART Algorithm

- 1 **Grow** a full tree T_0 on the training data using recursive binary splitting. Stopping criteria is *minsplit* which is the minimum number of observations that must exist in a node in order for a split to be attempted.
- 2 **Prune** the full tree T_0 to the subtree T_α using cost-complexity pruning.
- 3 **Determine** α using K-fold cross-validation, i.e. selecting best cp which is numerical value of α . In detail, divide the training set into K folds. For each $k = 1, \dots, K$:
 - 4 (a) Repeat Step 1 and 2 on all except for kth fold.
 - 5 (b) Calculate the mean squared prediction error on the hold out kth fold using T_α .
- 6 Finally, average the results from (b) for each value of α , and pick α to minimized the average prediction error.
- 7 **Return** the best subtree T_α .

Algorithm 1: CART: R-package-rpart

Regression Tree

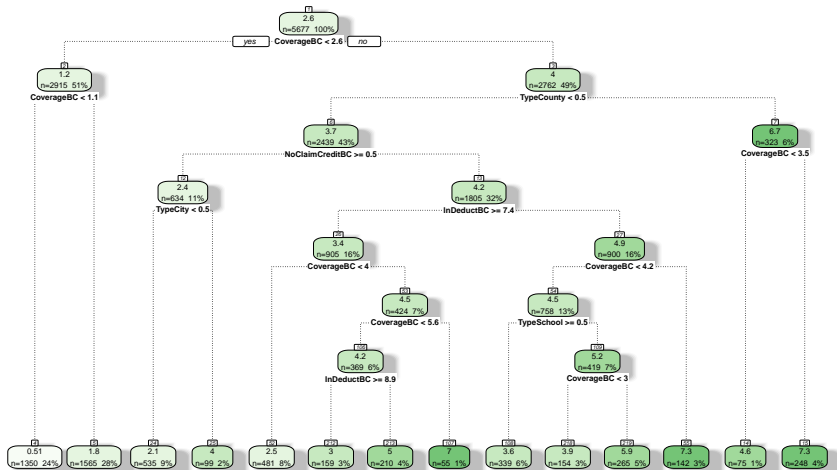


Figure 1: Univariate Regression Tree

Random Forest

- Random forest weakens the dependence among the CART trees by using a random subset of explanatory variables for split selection at each node of a tree.

- Ensemble methods

$$F_B(\mathbf{x}) = \frac{1}{B} \sum_{b=1}^B T(\mathbf{x}; \Theta_b)$$

- Train on bootstrap samplings B
- Random set of explanatory variables
- Decreases the correlation between different trained trees
- As $B \rightarrow \infty$, by the Strong Law of Large Numbers,

$$E_{\mathbf{x}, \mathbf{y}}(\mathbf{y} - F_B(\mathbf{x}))^2 \rightarrow E_{\mathbf{x}, \mathbf{y}}(\mathbf{y} - E_{\theta} T(\mathbf{x}; \Theta))^2 \quad a.s.$$

Random Forest Algorithm

- 1 **Bootstrap** For $b = 1$ to $B(ntree)$:
- 2 (a) Draw a bootstrap sample of size *sampsiz*e from the training data.
- 3 (b) Grow a full tree $T_b(\mathbf{x}; \Theta_b)$ on the bootstrap sample using recursive binary splitting. In detail, which different from Algorithm 1, select *mtry* variables at random from the p variables and stopping criteria is *nodesize*.
- 4 **Return** the ensemble of trees $\{T(\mathbf{x}; \Theta_b), b = 1, 2, \dots, B\}$.
- 5 **Averaging** $F_B(\mathbf{x}) = \frac{1}{B} \sum_{b=1}^B T(\mathbf{x}; \Theta_b)$

Algorithm 2: randomForest: R-package-randomForest

Bagging

- Bagging uses an ensemble of unpruned CART trees constructed from bootstrap samples of the data.

Gradient Boosted Regression Trees

- Boosting sequentially constructs the trees in the ensemble by putting more weight on the observations residuals in the previous step.

Gradient Boosted Regression Trees

- Builds regression trees sequentially on residuals

$$F_b(\mathbf{x}) = F_{b-1}(\mathbf{x}) + \sum_{m=1}^{M_b} c_{mb} \mathbf{1}_{R_{mb}}(\mathbf{x})$$

$$\hat{c}_{mb} = \arg \min_c \sum_{\mathbf{x}_i \in R_{mb}} L(y_i, F_{b-1}(\mathbf{x}_i) + c)$$

$$F_B(\mathbf{x}) = \sum_{b=1}^B T_b(\mathbf{x}; \Theta_b)$$

- The tree predictions $T_b(\mathbf{x}; \Theta_b)$ at each steps are analogous to the components of the negative gradient

$$g_{ib} = -\nabla_{F_{b-1}} L(y_i, F_{b-1}(\mathbf{x}_i)) = - \left[\frac{\partial L(y_i, F(\mathbf{x}_i))}{\partial F(\mathbf{x}_i)} \right]_{F(\mathbf{x}_i)=F_{b-1}(\mathbf{x}_i)}$$

Gradient Boosted Regression Trees Algorithm

- 1 **Initialize** model with a constant value, $F_0(x) = \arg \min_{\gamma} \sum_{i=1}^N L(y_i, \gamma)$.
- 2 **Gradient line search** For $b= 1$ to $B(n.trees)$:
 - 3 (a) For $i= 1$ to $N * p$, p is *bag.fraction*, compute $g_{ib} = -\nabla_{F_{b-1}} L(y_i, F_{b-1}(\mathbf{x}_i))$
 - 4 (b) Fit a regression tree $T_b(\mathbf{x}; \Theta_b)$ to the targets g_{ib} giving terminal regions R_1, R_2, \dots, R_{M_b} . In detail, which different from Algorithm 1, set *interaction.depth* = M_b and stopping criteria is *n.minobsinnode*.
 - 5 (c) For $m= 1$ to M_b compute $\hat{c}_{mb} = \arg \min_c \sum_{\mathbf{x}_i \in R_{mb}} L(y_i, F_{b-1}(\mathbf{x}_i) + c)$
 - 6 (d) Update $F_b(\mathbf{x}) = F_{b-1}(\mathbf{x}) + \lambda \sum_{m=1}^{M_b} c_{mb} \mathbf{1}_{R_{mb}}(\mathbf{x})$ here we can add *shrinkage* λ to reduce the impact of each additional fitted base-learner, regression tree, $T_b(\mathbf{x}; \Theta_b)$.
- 7 **Return** $F_B(\mathbf{x}) = \sum_{b=1}^B \lambda T_b(\mathbf{x}; \Theta_b)$.

Algorithm 3: gbm: R-package-gbm

Multivariate regression trees

- Extended the univariate regression trees
- Multivariate measure

$$L(\mathbf{y}_i, \hat{\mathbf{y}}_i) = \sum_{i=1}^N (\mathbf{y}_i - \hat{\mathbf{y}}_i)^T (\mathbf{y}_i - \hat{\mathbf{y}}_i)$$

$$L(\mathbf{y}_i, \hat{\mathbf{y}}_i) = \sum_{i=1}^N \sum_{j=1}^k |y_{ij} - \tilde{y}_j|$$

- Inherited univariate decision tree advantages
- Robust to the outliers and missing data
- Constrained clustering